




Host DRAM — overflow only (priority 3)

reached only when a block exceeds a GPU slab, or every slab is live

serving
 waiting — session paused between turns
 completed
 evicted / absent

the cache cannot tell waiting from completed — both are LRU-evictable; that is what parking is for

- 1 R_{10} arrives; the local radix misses — its prefix was evicted under pressure
- 2 the park index finds $KV_{1..3}$ on the idle decode GPU
- 3 peer-GPU copy back into the prefill KV pool and radix
- 4 prefill only the new tokens; re-park the extended prefix